

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\6_Salesworkload1.csv")
df.fillna(0,inplace=True)
df
```

Out[2]:

	MonthYear	Time index	Country	StoreID	City	Dept_ID	Dept. Name	HoursOwn	HoursLease	Sa ui
0	10.2016	1.0	United Kingdom	88253.0	London (I)	1.0	Dry	3184.764	0.0	39856
1	10.2016	1.0	United Kingdom	88253.0	London (I)	2.0	Frozen	1582.941	0.0	8272
2	10.2016	1.0	United Kingdom	88253.0	London (I)	3.0	other	47.205	0.0	43840
3	10.2016	1.0	United Kingdom	88253.0	London (I)	4.0	Fish	1623.852	0.0	30942
4	10.2016	1.0	United Kingdom	88253.0	London (I)	5.0	Fruits & Vegetables	1759.173	0.0	16551
...
7653	06.2017	9.0	Sweden	29650.0	Gothenburg	12.0	Checkout	6322.323	0.0	388653
7654	06.2017	9.0	Sweden	29650.0	Gothenburg	16.0	Customer Services	4270.479	0.0	24
7655	06.2017	9.0	Sweden	29650.0	Gothenburg	11.0	Delivery	0	0.0	
7656	06.2017	9.0	Sweden	29650.0	Gothenburg	17.0	others	2224.929	0.0	24
7657	06.2017	9.0	Sweden	29650.0	Gothenburg	18.0	all	39652.2	0.0	388653

7658 rows × 14 columns

```
In [3]: df.head()
```

Out[3]:

	MonthYear	Time index	Country	StoreID	City	Dept_ID	Dept. Name	HoursOwn	HoursLease	Sales units	Tu
0	10.2016	1.0	United Kingdom	88253.0	London (I)	1.0	Dry	3184.764	0.0	398560.0	122
1	10.2016	1.0	United Kingdom	88253.0	London (I)	2.0	Frozen	1582.941	0.0	82725.0	38
2	10.2016	1.0	United Kingdom	88253.0	London (I)	3.0	other	47.205	0.0	438400.0	65
3	10.2016	1.0	United Kingdom	88253.0	London (I)	4.0	Fish	1623.852	0.0	309425.0	49
4	10.2016	1.0	United Kingdom	88253.0	London (I)	5.0	Fruits & Vegetables	1759.173	0.0	165515.0	32

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7658 entries, 0 to 7657
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   MonthYear              7658 non-null   object
1   Time index             7658 non-null   float64
2   Country                7658 non-null   object
3   StoreID                7658 non-null   float64
4   City                   7658 non-null   object
5   Dept_ID                7658 non-null   float64
6   Dept. Name             7658 non-null   object
7   HoursOwn               7658 non-null   object
8   HoursLease             7658 non-null   float64
9   Sales units            7658 non-null   float64
10  Turnover               7658 non-null   float64
11  Customer               7658 non-null   float64
12  Area (m2)              7658 non-null   object
13  Opening hours          7658 non-null   object
dtypes: float64(7), object(7)
memory usage: 837.7+ KB
```

```
In [5]: import seaborn as sns
```

```
In [6]: df.describe()
```

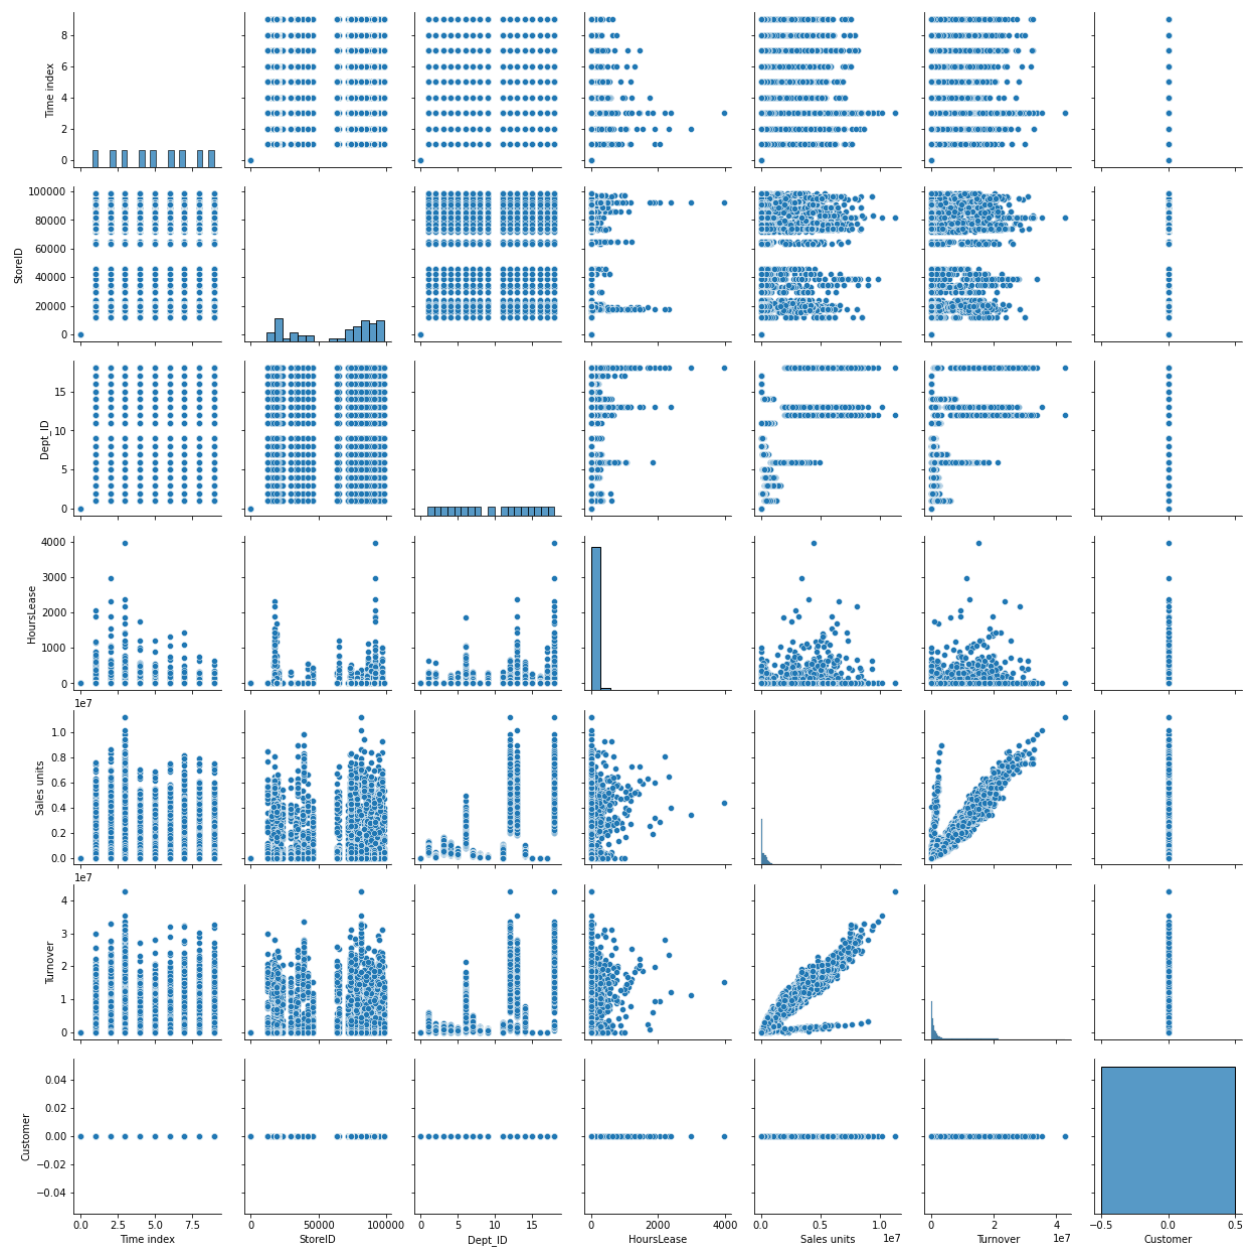
Out[6]:

	Time index	StoreID	Dept_ID	HoursLease	Sales units	Turnover	Customer
count	7658.000000	7658.000000	7658.000000	7658.000000	7.658000e+03	7.658000e+03	7658.0
mean	4.994777	61930.456124	9.460695	22.013058	1.075346e+06	3.717505e+06	0.0
std	2.585859	29975.929873	5.343407	133.231761	1.727560e+06	6.001448e+06	0.0
min	0.000000	0.000000	0.000000	0.000000	0.000000e+00	0.000000e+00	0.0
25%	3.000000	29650.000000	5.000000	0.000000	5.441375e+04	2.720558e+05	0.0
50%	5.000000	73949.000000	9.000000	0.000000	2.927625e+05	9.300810e+05	0.0
75%	7.000000	87703.000000	14.000000	0.000000	9.154812e+05	3.251488e+06	0.0
max	9.000000	98422.000000	18.000000	3984.000000	1.124296e+07	4.271739e+07	0.0

```
In [ ]:
```

```
In [7]: sns.pairplot(df)
```

Out[7]: <seaborn.axisgrid.PairGrid at 0x1921cc61be0>




```
In [8]: df1=df.drop(['Country'],axis=1)
df1
df1=df1.drop(df1.index[1537:])
df1.isna().sum()
```

```
Out[8]: MonthYear      0
Time index      0
StoreID         0
City            0
Dept_ID         0
Dept. Name      0
HoursOwn        0
HoursLease      0
Sales units     0
Turnover        0
Customer        0
Area (m2)       0
Opening hours   0
dtype: int64
```


```
In [9]: sns.displot(df['StoreID'])
```

```
Out[9]: <seaborn.axisgrid.FacetGrid at 0x19219e54610>
```



```
In [10]: sns.heatmap(df1.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression
```

```
In [12]: df1.isna().sum()
```

```
Out[12]: MonthYear      0  
Time index      0  
StoreID         0  
City           0  
Dept_ID        0  
Dept. Name     0  
HoursOwn       0  
HoursLease     0  
Sales units    0  
Turnover       0  
Customer       0  
Area (m2)      0  
Opening hours  0  
dtype: int64
```

```
In [13]: y=df1['Turnover']
x=df1.drop(['Turnover','MonthYear','City','Opening hours','Dept. Name','Customer','Turno
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
print(x_train)
```

	Time index	StoreID	Dept_ID	HoursOwn	HoursLease	Sales units
1147	2.0	85696.0	7.0	5680.335	84.0	240170.0
572	1.0	12227.0	15.0	3974.661	0.0	165.0
1515	2.0	79785.0	2.0	2278.428	0.0	92455.0
1460	2.0	34378.0	11.0	0	0.0	10.0
1217	2.0	19000.0	9.0	1595.529	0.0	41060.0
...
662	1.0	98422.0	18.0	46276.635	0.0	3508675.0
783	1.0	73762.0	2.0	2426.337	0.0	182250.0
1246	2.0	20166.0	5.0	1992.051	0.0	243510.0
1242	2.0	20166.0	1.0	3449.112	0.0	557495.0
542	1.0	78325.0	17.0	2193.459	0.0	155.0


[1075 rows x 6 columns]

```
In [14]: model=LinearRegression()
model.fit(x_train,y_train)
model.intercept_
```

Out[14]: 71487.15112321172

```
In [15]: prediction=model.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[15]: <matplotlib.collections.PathCollection at 0x1922262a670>



```
In [16]: model.score(x_test,y_test)
```

Out[16]: 0.8844082866440194

```
In [17]: from sklearn.linear_model import Ridge,Lasso
```

```
In [18]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[18]: Ridge(alpha=10)
```

```
In [19]: rr.score(x_test,y_test)
```

```
Out[19]: 0.8844074382998153
```

```
In [20]: la =Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[20]: Lasso(alpha=10)
```

```
In [21]: la.score(x_test,y_test)
```

```
Out[21]: 0.884408249300056
```

```
In [22]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[22]: ElasticNet()
```

```
In [23]: print(en.coef_)
```

```
[-8.49670897e+03 -2.83657677e+00  2.80818729e+04  2.60727171e+01
  3.91533921e+01  3.07359204e+00]
```

```
In [24]: print(en.intercept_)
```

```
51211.900752046145
```

```
In [25]: print(en.predict(x_test))
```

```
51211.900752046145 51211.900752046145 51211.900752046145 51211.900752046145
12081313.10615316  318344.55719734  161808.63620064 20386487.79747214
1104261.62469585  1985757.29088112  1515331.45219829  154903.21553394
1250542.62909192  546798.39424928  866228.38947379  1267149.35330232
1049485.95428433  17454798.87082483  18277896.92037161  526152.89263141
322212.53734578  298812.19160697  1101221.19750973  245088.25842468
341348.9338572  1142780.75439107  335986.78491951  487380.15303624
2018371.9873523  532543.62839248  1670797.07024807  314843.80322381
1946713.52384728  526580.01513346  4691939.55389973  407220.35528643
527200.7612785  7427458.75733612  1590263.56950927  736289.56119032
1815862.82154166  18885560.24321736  383570.58251211  529899.46255776
415955.72426126  314857.84317711  360335.27297  1483195.17038797
2833594.396895  1772604.38529096  269436.8816465  15378673.112774
8980402.98260365  422353.62165203  120944.20925279  142384.21800239
171731.73735571  7451817.89593718  466892.72012039  336644.33983013
6034595.82323731  11392712.30266854  16354156.15313049  1533643.13924037
870754.86557243  5316510.32895393  8452937.38126709  277696.84394422
1007142.9373982  292356.86817356  395279.80168817  1120197.39862718
937535.45049221  1330982.36569508  246103.72548528  7537691.05049852
3879704.8851028  400243.89950139]
```

```
In [26]: print(en.score(x_test,y_test))
```

```
0.8844007933875692
```

```
In [27]: from sklearn import metrics
```

```
In [28]: print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolute Error: 1003905.7077692078
```

```
In [29]: print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

```
Mean Squared Error: 4032719243820.1797
```

```
In [30]: print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction))
```

```
Root Mean Squared Error: 2008163.1516936515
```