

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\11_winequality-red.csv")
df.fillna(0,inplace=True)
df
```

Out[2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	
...	
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	1
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	1
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	1
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	1
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	1

1599 rows × 12 columns



```
In [3]: df.head()
```

Out[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4



In [4]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density               1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates             1599 non-null   float64
10  alcohol               1599 non-null   float64
11  quality               1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

In [5]: import seaborn as sns

In [6]: df.describe()

Out[6]:

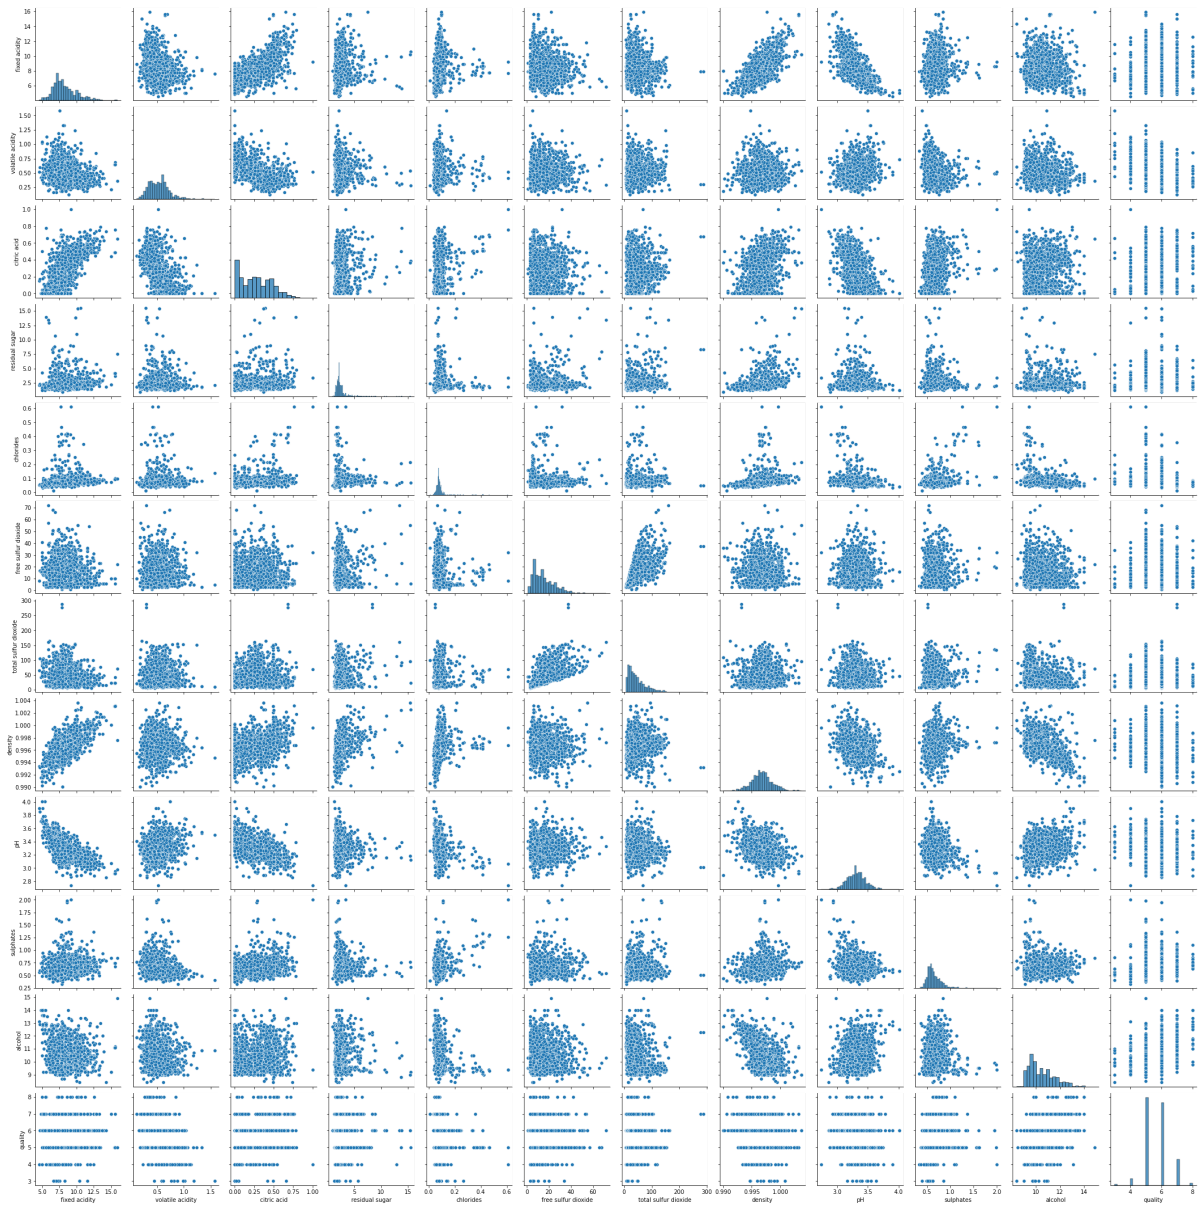
	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.46779
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.89532
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.00000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.00000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.00000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.00000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.00000



In []:

```
In [7]: sns.pairplot(df)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x1fdf66b4e80>
```




```
In [8]: df1=df.drop(['citric acid'],axis=1)
df1
df1=df1.drop(df1.index[1537:])
df1.isna().sum()
```

```
Out[8]: fixed acidity      0
volatile acidity    0
residual sugar      0
chlorides           0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                0
sulphates          0
alcohol            0
quality            0
dtype: int64
```

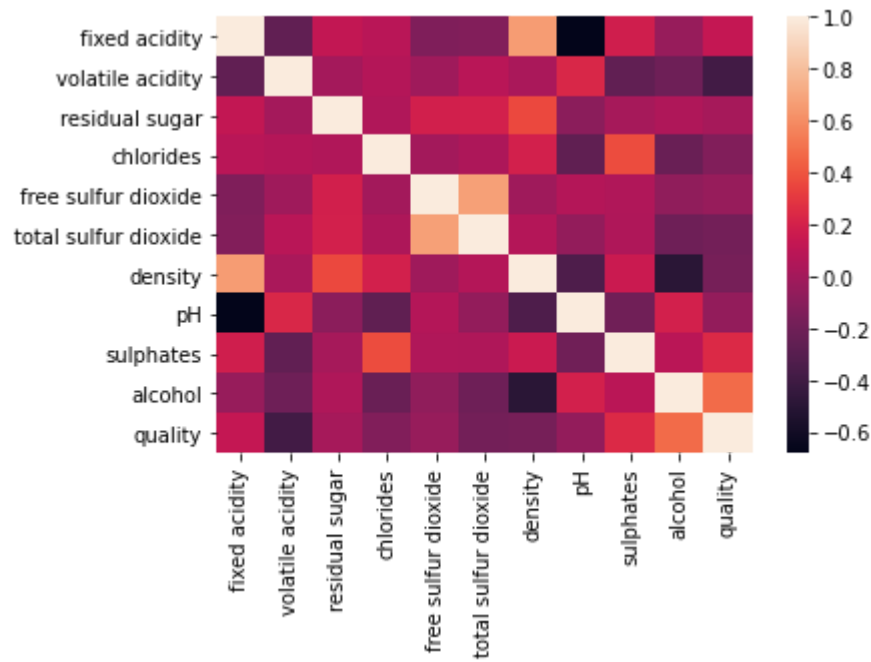
```
In [9]: sns.displot(df['alcohol'])
```

```
Out[9]: <seaborn.axisgrid.FacetGrid at 0x1fd6da65190>
```



```
In [10]: sns.heatmap(df1.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression
```

```
In [12]: df1.isna().sum()
```

```
Out[12]: fixed acidity      0  
volatile acidity    0  
residual sugar      0  
chlorides           0  
free sulfur dioxide  0  
total sulfur dioxide 0  
density            0  
pH                 0  
sulphates          0  
alcohol            0  
quality            0  
dtype: int64
```

```
In [13]: y=df1['fixed acidity']
x=df1.drop(['chlorides','residual sugar'],axis=1)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
print(x_train)
```

	fixed acidity	volatile acidity	free sulfur dioxide \
1341	7.5	0.51	13.0
801	8.6	0.55	8.0
1459	7.9	0.20	7.0
756	6.3	0.98	15.0
1126	5.8	0.29	3.0
...
583	12.0	0.28	10.0
1518	7.4	0.47	7.0
1490	7.1	0.22	8.0
388	7.8	0.46	23.0
118	8.8	0.55	14.0

	total sulfur dioxide	density	pH	sulphates	alcohol	quality
1341	31.0	0.99538	3.36	0.54	10.5	6
801	17.0	0.99735	3.23	0.44	10.0	5
1459	15.0	0.99458	3.32	0.80	11.9	7
756	33.0	0.99488	3.60	0.46	11.2	6
1126	11.0	0.99150	3.39	0.54	13.5	6
...
583	21.0	0.99760	2.98	0.66	9.9	7
1518	20.0	0.99647	3.32	0.63	10.5	5
1490	18.0	0.99344	3.39	0.56	12.4	6
388	53.0	0.99810	3.43	0.74	9.2	6
118	56.0	0.99620	3.21	0.60	10.9	6


[1075 rows x 9 columns]

```
In [14]: model=LinearRegression()
model.fit(x_train,y_train)
model.intercept_
```

Out[14]: 7.993605777301127e-14

```
In [15]: prediction=model.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x1fdfffd7610>
```



```
In [16]: model.score(x_test,y_test)
```

```
Out[16]: 1.0
```

```
In [17]: from sklearn.linear_model import Ridge,Lasso
```

```
In [18]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[18]: Ridge(alpha=10)
```

```
In [19]: rr.score(x_test,y_test)
```

```
Out[19]: 0.9999872504686617
```

```
In [20]: la =Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[20]: Lasso(alpha=10)
```

```
In [21]: la.score(x_test,y_test)
```

```
Out[21]: -0.00024395652394071377
```

```

In [22]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)

print(en.coef_)

print(en.intercept_)

print(en.predict(x_test))

print(en.score(x_test,y_test))

from sklearn import metrics

print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))

print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))

print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))

```

10.32023011	10.07073307	8.10000273	8.30434330	8.00324371	7.03077030
7.67156836	7.04546752	9.849262	10.04525574	7.43595514	8.32032773
7.95160255	8.05184922	13.49796054	12.15211542	8.15569941	8.29091502
6.79844956	7.91693551	8.89210275	7.46386799	8.26615284	9.38628961
8.50176607	7.5987817	9.98762828	6.78689388	8.10302439	10.63173711
9.64336338	8.78975242	9.25617738	8.00382474	8.40947156	7.09334105
10.72238081	8.33203436	7.64170281	9.4424172	5.98460902	8.92016654
6.25308753	7.04546752	7.930142	7.49508247	5.98610889	8.23328756
7.06197563	7.80813289	7.81158546	7.55451168	11.53937202	9.86907174
8.63517992	8.15554846	11.6496745	8.08321465	10.49832326	9.7224514
7.59248035	7.37157349	8.71426793	8.59075895	8.80326081	8.78840351
10.71247594	8.33203436	7.29908872	6.9663795	11.48339537	9.05673107
8.21993012	8.51497256	8.8854995	7.15096851	7.44090757	10.34675049
8.78840351	8.66639439	7.17557974	9.93165163	11.05493908	8.58085408
10.95619227	6.46408953	6.46739115	7.97621378	7.4077404	10.18692365
7.22990558	5.92037831	7.86921292	7.7949264	9.20185155	7.70608446]

0.9155323766730585
 Mean Absolute Error: 7.782134724991574e-15
 Mean Squared Error: 9.409300658681618e-29
 Root Mean Squared Error: 9.700154977463823e-15